

Computer Vision

Exercises of Lab 6

Exercise 6.1: Optical flow

The goal of this exercise is to understand the principles of optical flow and experiment with motion models such as pure translation and affine motion.

Make sure that you have VLFeat. Open Exercise6.1.m and change the variable VLFEATROOT to the unpacked paths at your hard drive.

Look through the code where the strongest Harris corners features are detected in frame 1. We want to track these points in the subsequent frames using optical flow. When the motion between subsequent frames is small, this is a more computationally cheap approach than e.g. estimating SIFT features and matching them between frames.

Construct the matrices M and b used for estimating optical flow. Look at slide 37-41 in Lecture5_OpticalFlow.pdf for reference. Use the variables I_x , I_y and I_t defined in the code.

For tracking the feature points, a pure translational motion model is used. However, for measuring the similarity of a feature from frame 1 to i , it is better to use an affine motion model. Why would an affine transformation map a feature window better than a pure translation?

Estimate an affine transformation between frame 1 and i for a feature point j . That is, construct the matrices A and b that define a linear least squares problem of estimating the affine motion. Look at slide 56 for reference. Here, we want to estimate the parameters a_1 - a_6 . We want to estimate the vector x in the linear equations: $Ax=b$ The solution to such a problem that minimizes $\|Ax-b\|^2$ is given by:

$$x^*=(A^TA)^{-1}A^Tb$$

where x^* consists of the parameters a_1 - a_6 .

Exercise 6.2: Canny edge detection

The goal of this exercise is to understand the Canny edge detector and implement non-maximum suppression.

Run Exercise6.2.m, read the code and inspect Figure 1 carefully.

We are going to implement non-maximum suppression, constructing thin edge lines in the output image. To do this, construct two gradient vectors (opposite directions) with length 1 for the point (i,j) . The x - and y -values should be stored in the vectors, u and v , respectively. Look at slide 33 in Lecture5.1-2_EdgesLines.pdf.


Hint: use the gradient maps, I_x and I_y , and normalize with I_{norm} .

Then, plot the gradient vectors on top of the gradient map using the quiver function.

Finally, implement the actual edge thinning by assigning $EdgeMap(i,j)=1$ if the current gradient at (i,j) is larger than or equal to its two interpolated neighbors.

Exercise 6.3: Texture gradients

The goal of this exercise is to detect the edge between two textured areas.



Make sure that you have VLFeat. Open Exercise6.3.m and change the variable VLFEATROOT to the unpacked paths at your hard drive.

Run the Matlab script and study Figure 1 (the program pauses four times – hit any key to proceed). Notice how the differently oriented filters provide significantly different outputs.

Cluster the pixels in the image using the outputs of the 4 different filters. You should construct a variable, datapoints, with 4-dimensional feature vectors and feed them to the clustering algorithm kmeans. You can use VLFeat's implementation `vl_kmeans(datapoints,numClust)` that takes datapoints and clusters them into numClust clusters.

Following the clustering procedure, construct a texton map, textonMap. Each pixel should be assigned a cluster number as its intensity. Use `show_norm_image(textonMap)` to visualize it.

Run the remainder of the code and study it carefully. Make sure that you understand exactly what happens in Figure 3 and 4.

Try changing the different parameters of the algorithm and look at their influence. Change the number of clusters, numClust. Change the number and orientations of the directional filters.